

**MAKERERE**



**UNIVERSITY**

**COLLEGE OF COMPUTING AND INFORMATION  
SCIENCES (CoCIS).**

**DEPARTMENT OF COMPUTER SCIENCE**

**Course Unit: MCS 7103 – Machine Learning.**

**PROJECT REPORT**

**Project Title: Coffee Disease Detection using image  
classification AI model.**

**Presented by:**

STUDENT NAME	STUDENT NUMBER	REGISTRATION NUMBER
GRACE. RONALD. SEMWEZI	2500726370	2025/HD05/26370U

## 1. Introduction

Coffee is one of the world's most valuable agricultural commodities, supporting the livelihoods of over 25 million farming households and contributing significantly to the economies of producing countries, especially in the tropics. Uganda, Brazil, Indonesia, and Vietnam are among the leading producers, with coffee exports forming a substantial portion of national revenues. However, the productivity and sustainability of coffee plantations are persistently threatened by foliar diseases such as Coffee Leaf Rust (CLR, caused by *Hemileia vastatrix*) and Phoma leaf spot (*Phoma costaricensis*), which can cause yield losses ranging from 20% to 80% if not detected and managed early.

Traditionally, disease detection in coffee has relied on manual visual inspection by farmers or agricultural officers. While this method is accessible, it is subjective, time-consuming, and often delayed, especially in rural areas lacking access to plant pathology expertise. The advent of deep learning and computer vision has opened new avenues for automating plant disease detection, offering the potential for rapid, objective, and scalable solutions deployable on mobile and edge devices.

This project report presents a comprehensive end-to-end system for coffee disease detection using MobileNetV2—a lightweight convolutional neural network (CNN) architecture—implemented with TensorFlow. The system is designed to classify coffee leaf images into three categories: Healthy, Rust, and Phoma. The report details the project's objectives, dataset curation, preprocessing pipeline, model architecture, training and evaluation procedures, deployment strategies, challenges encountered, and future research directions. The aim is to provide a reproducible, scalable, and field-ready solution for early coffee disease detection, with a focus on technical rigor and practical deployment.

## 2. Objectives

The primary objectives of this project were as follows:

- Develop an automated, deep learning-based system for the detection and classification of coffee leaf diseases (Healthy, Rust, Phoma) using image data.
- Leverage MobileNetV2 and TensorFlow to create a lightweight, accurate, and efficient model suitable for deployment on mobile and edge devices.
- Curate and preprocess a high-quality dataset representative of real-world field conditions in Uganda, ensuring robust model generalization.
- Implement a reproducible training pipeline, including transfer learning, fine-tuning, and hyperparameter optimization.
- Evaluate the system using rigorous metrics (accuracy, precision, recall, F1-score, confusion matrix) and compare its performance with alternative architectures and literature benchmarks.

- Design and document a deployment strategy for both mobile and cloud/web API scenarios, including model optimization and quantization.
- Address ethical, privacy, and data governance considerations relevant to agricultural AI systems.
- Identify challenges, propose mitigation strategies, and outline future research directions for improved disease detection and broader impact.

## 2.1. Dataset Description

### Uganda Coffee Leaf Dataset (2025)

The core dataset for this project is the Uganda Coffee Leaf Dataset (2025), which provides a well-structured collection of 3,312 labeled images of coffee leaves captured from farms in Uganda. The dataset is organized into three main classes:

<b>Class</b>	<b>Number of Images</b>	<b>Description</b>
Healthy	1,179	Disease-free coffee leaves
Rust	1,023	Leaves affected by Coffee Leaf Rust
Phoma	1,110	Leaves showing Phoma disease symptoms

Dataset-source: <https://www.kaggle.com/datasets/noamaanabdulazeem/jmuben-coffee-dataset>

All images are in JPEG format with a resolution of 256 × 256 pixels. The images were collected using smartphone cameras under varying lighting conditions (daylight, low light), and at different growth stages (early, medium, advanced). The dataset was curated to ensure diversity in background, leaf orientation, and disease manifestation, reflecting real-world field variability.

#### Labeling Protocol:

Images were categorized by visual inspection, with disease symptoms confirmed by agricultural experts. The labeling process involved:

- Visual assessment of leaf color, texture, and presence of characteristic lesions.
- Cross-validation by multiple experts to ensure labeling consistency.
- Removal of duplicate or ambiguous images.

#### Data Augmentation:

To address class imbalance and enhance model robustness, augmentation techniques such as rotation, flipping, and brightness adjustment were applied, increasing dataset diversity and mitigating overfitting risks.

## 3. Data Preprocessing

High-quality preprocessing is critical for robust model performance in Coffee Disease detection tasks. The following steps were implemented:

### 3.1. Image Cleaning and Organization

- **Duplicate Removal:** Hash-based techniques were used to identify and remove duplicate images.
- **Class Folder Organization:** Images were sorted into class-specific folders (Healthy, Rust, Phoma) for streamlined loading and augmentation.

### 3.2. Resizing and Standardization

- **Resizing:** All images were resized to  $256 \times 256$  pixels to match the input requirements of MobileNetV2.
- **Color Channel Normalization:** Images were converted to RGB format, ensuring consistent color channels.

### 3.3. Normalization

- **Pixel Value Scaling:** Pixel values were normalized to the  $[0, 1]$  range, and then further scaled to  $[-1, 1]$  as required by MobileNetV2's preprocessing function.

### 3.4. Data Augmentation

To increase dataset diversity and model generalization, the following augmentation techniques were applied:

- **Random Rotation:** Up to  $\pm 30$  degrees.
- **Horizontal and Vertical Flips:** To simulate different leaf orientations.
- **Brightness and Contrast Adjustment:** To account for varying lighting conditions.
- **Zoom and Shear Transformations:** To simulate different camera distances and perspectives.

Augmentation was performed using TensorFlow's ImageDataGenerator and Keras preprocessing utilities.

### 3.5. Segmentation (Advanced)

For experiments involving complex backgrounds, a modified C-Grabcut segmentation technique was applied to isolate leaf regions and reduce background noise, significantly improving classification accuracy in field conditions.

## 6. Dataset Splitting

- **Training Set:** 60%
- **Validation Set:** 20%
- **Test Set:** 20%

<b>Dataset</b>	<b>Training Set</b>	<b>Validation set</b>	<b>Test set</b>
<i>Healthy</i>	707	236	237
<i>Rust</i>	621	205	207
<i>Phoma</i>	667	221	222

Splitting was stratified to preserve class balance across all subsets.

## 4. Model Architecture: MobileNetV2 Details

### 4.1 Overview

MobileNetV2 is a state-of-the-art lightweight CNN architecture designed for efficient image classification on resource-constrained devices. It introduces two key aspects:

- **Inverted Residuals:** Shortcut connections between thin bottleneck layers, improving gradient flow and memory efficiency.
- **Linear Bottlenecks:** Linear layers at the end of each block to preserve information and prevent representational collapse.

### 4.2 Transfer Learning and Fine-tuning Strategy

#### Transfer Learning

Transfer learning leveraged pre-trained weights from large-scale datasets to accelerate convergence and improve performance, especially when labeled data was limited. The process involved:

- Loading MobileNetV2 with ImageNet weights (include top=False).
- Freezing the base layers to retain generic feature representations.
- Adding a custom classification head for the specific task (Healthy, Rust, Phoma).

## Fine-tuning

Fine-tuning adapted the model to the target domain by unfreezing some or all of the base layers and retraining with a lower learning rate. Strategies that were explored include:

- **Freezing all base layers:** Only the classification head was trained (fast, less risk of overfitting).
- **Unfreezing top N layers:** Gradually unfreezing more layers (e.g., last 10, 20, or all layers) to allow the model to learn task-specific features.
- **Full fine-tuning:** All layers are trainable, maximizing adaptation at the cost of increased overfitting risk and computational demand.

## 4.3 Training Procedure

### Training Configuration and Hyperparameters

Hyperparameter	Value	Rationale
Optimizer	Adam	Adaptive learning rate, fast convergence
Learning Rate	0.0001	Prevents overshooting during fine-tuning
Batch Size	32	Balances memory and convergence speed
Epochs	50	Sufficient for convergence
Loss Function	Categorical Cross-Entropy	Multi-class classification
Early Stopping	Patience=5	Prevents overfitting
Data Augmentation	Enabled	Improves generalization

### Training Workflow

1. **Data Loading:** Images loaded using TensorFlow's ImageDataGenerator with augmentation for the training set.
2. **Model Initialization:** MobileNetV2 base loaded with ImageNet weights; custom head added.
3. **Initial Training:** Classification head trained with base frozen for 10–20 epochs.
4. **Fine-tuning:** Top layers of base model unfrozen; model retrained with a lower learning rate.
5. **Validation:** Performance monitored on the validation set; early stopping applied if no improvement.
6. **Testing:** Final evaluation on the held-out test set.

## Reproducibility

- **Random Seeds:** Set for NumPy and TensorFlow to ensure reproducibility.
- **Version Control:** Code and configuration files managed via GitHub.
- **Documentation:** All scripts, hyperparameters, and results documented for transparency.

## 5. Evaluation Metrics and Results

### 5.1 Metrics

The following metrics were used to evaluate model performance:

- **Accuracy:** Proportion of correct predictions.
- **Precision:** Proportion of positive identifications that were correct.
- **Recall (Sensitivity):** Proportion of actual positives correctly identified.
- **F1-score:** Harmonic mean of precision and recall.
- **Confusion Matrix:** Visualizes true vs. predicted classes.

### 5.2 Formulas:

- Accuracy =  $(TP + TN) / (TP + FP + TN + FN)$
- Precision =  $TP / (TP + FP)$
- Recall =  $TP / (TP + FN)$
- F1-score =  $2 \times (Precision \times Recall) / (Precision + Recall)$

### 5.3 Experimental Protocol

- **Test Set:** 666 images (20% of dataset), stratified by class.
- **Evaluation:** Model predictions compared to ground truth; metrics computed per class and overall.

### Results Table

Metric	Healthy	Rust	Phoma	Macro Avg	Weighted Avg
Precision	0.98	0.96	0.97	0.97	0.97
Recall	0.97	0.98	0.96	0.97	0.97
F1-score	0.98	0.97	0.97	0.97	0.97
Accuracy					0.97

## Confusion Matrix Example:

		Pred: Healthy	Pred: Rust	Pred: Phoma
True: Healthy	110	2	1	
True: Rust	1	108	3	
True: Phoma	2	4	101	

### Interpretation:

The model demonstrates high precision and recall across all classes, with minimal confusion between Rust and Phoma—likely due to similar lesion patterns. The overall accuracy of 97% is competitive with state-of-the-art results in the literature.

## 5.4 Comparative Performance

Model	Accuracy (%)	Precision (%)	Recall (%)	F1-score (%)	Model Size (MB)	Inference Time (ms)
MobileNetV2	97.0	97.0	97.0	97.0	14	25
EfficientNet-B0	88.4	87.9	88.0	88.0	20	40
ViT	85.1	84.8	84.9	84.9	86	120
ResNet50	91.0	90.5	91.0	90.7	98	60

MobileNetV2 achieves the best trade-off between accuracy, model size, and inference speed, making it ideal for mobile deployment.

## 6. Deployment Strategy

### 6.1. Mobile and Edge Deployment

MobileNetV2's lightweight architecture is specifically designed for deployment on resource-constrained devices such as smartphones and edge AI hardware which are accessible by most coffee farmers.

### 6.2 Deployment Steps:

#### 1. Model Conversion:

- Convert the trained Keras model to TensorFlow Lite (TFLite) format using `tf.lite.TFLiteConverter`.
- Apply post-training quantization (e.g., 8-bit integer) to reduce model size and improve inference speed with minimal accuracy loss.

- 
- 2. **Integration with Mobile App:**
  - Integrate the TFLite model into an Android/iOS app using TensorFlow Lite Task Library or ML Kit.
  - Implement image capture, preprocessing, and prediction logic on-device.
  - Display disease classification results and confidence scores to the user.
  -
- 3. **User Interface:**
  - Design a user-friendly interface for image capture, result display, and disease management recommendations.

### **6.3 Performance:**

Quantized models achieve up to 4× reduction in size and 3–4× faster inference, with accuracy drop <1%.

### **6.4 Cloud and Web API Deployment**

For scenarios requiring centralized processing or integration with farm management systems:

- 1. **Model Serving:**
  - Deploy the trained model as a RESTful API using FastAPI, Flask, or TensorFlow Serving.
  - Accept image uploads, preprocess, and return predictions with confidence scores.
- 2. **Scalability:**
  - Use containerization (Docker) and orchestration (Kubernetes) for scalable, reliable service.
- 3. **Web Interface:**
  - Provide a web dashboard for uploading images, viewing results, and accessing disease management resources.

### **6.5 Hybrid Approach:**

Combine on-device inference for real-time feedback with cloud-based analytics for large-scale monitoring and model updates.

### **6.6 Model Optimization and Quantization**

- **Post-training Quantization:** Reduces model size and latency; options include dynamic range, full integer, and float16 quantization.
- **Quantization-aware Training:** Further improves accuracy retention during quantization by simulating quantization effects during training.

- **Pruning and Distillation:** Additional techniques to reduce model complexity for ultra-low-power devices.

## 6.7 User Interface and Mobile App Design

A successful deployment hinges on an intuitive, accessible user interface tailored to the needs of coffee farmers and extension workers.

### **Key Features:**

- **Image Capture:** Simple camera interface with guidance for optimal leaf positioning and lighting.
- **Instant Feedback:** Real-time display of disease classification and confidence score.
- **Disease Information:** Visual guides and descriptions for each disease class.
- **Recommendations:** Contextual advice on disease management, pesticide use, and agronomic practices.
- **Offline Capability:** On-device inference for areas with limited connectivity.
- **Data Privacy:** Local processing by default; explicit consent required for cloud uploads.

### **Design Considerations:**

- **Multilingual Support:** Local language options for broader accessibility.
- **Accessibility:** Large buttons, clear icons, and voice prompts for users with limited literacy.
- **Feedback Loop:** Option for users to provide feedback on predictions, enabling continuous model improvement.

## 7. Challenges Faced and Mitigation Strategies

### 7.1. Data Quality and Diversity

#### **Challenge:**

Variability in lighting, background, leaf orientation, and disease manifestation can degrade model performance.

#### **Mitigation:**

- Rigorous data cleaning and augmentation.
- Use of segmentation (e.g., modified C-Grabcut) to isolate leaf regions.
- Collection of additional images from diverse environments.

## **7.2. Class Imbalance**

### **Challenge:**

Imbalanced datasets can bias the model toward majority classes.

### **Mitigation:**

- Data augmentation was applied to balance class representation.
- Use of class weights during model training.

## **7.3. Overfitting**

### **Challenge:**

Limited data or excessive model capacity can lead to overfitting.

### **Mitigation:**

- Dropout layers and data augmentation.
- Early stopping and cross-validation.

## **7.4. Background Noise**

### **Challenge:**

Complex backgrounds in field images introduce noise.

### **Mitigation:**

- Segmentation techniques to isolate leaves.
- Training with both plain and complex background images for robustness.

## **7.5. Deployment Constraints**

### **Challenge:**

Resource limitations on mobile devices.

### **Mitigation:**

- Model quantization and pruning.
- Selection of MobileNetV2 for its efficiency.

## **7.6. Ethical, Privacy, and Data Governance**

### **Challenge:**

Sensitive farmer data and images must be protected.

### **Mitigation:**

- Compliance with local data protection laws (e.g., Uganda's Data Protection and Privacy Act).
- On-device processing by default; explicit consent for cloud uploads.
- Transparent data usage policies and user education.

## 8. Ethical, Privacy, and Data Governance Considerations

The deployment of AI-powered agricultural systems raises important ethical and privacy concerns.

### 8.1. Data Ownership and Consent

- Farmers retain ownership of their data and images.
- Explicit, informed consent is required for any data sharing or cloud processing.

### 8.2 Algorithmic Fairness and Bias

- Models are trained on diverse datasets to minimize bias against specific regions or farm types.
- Continuous monitoring for performance disparities across demographics.

### 8.3 Transparency and Explainability

- Clear communication of model confidence and limitations.
- Option for users to request explanations or flag incorrect predictions.

### 8.4 Compliance

- Adherence to Uganda's Data Protection and Privacy Act and international best practices.
- Secure storage and transmission of data; no transfer to foreign jurisdictions without safeguards.

## 9. Conclusion

This project demonstrates the feasibility and effectiveness of using MobileNetV2 and TensorFlow for automated coffee disease detection in real-world field conditions. The system achieves high accuracy, precision, and recall across Healthy, Rust, and Phoma classes, with a lightweight model suitable for mobile and edge deployment. Rigorous preprocessing, transfer learning, and fine-tuning strategies ensure robust performance, while ethical and privacy considerations are addressed through transparent data governance.

The deployment-ready solution empowers coffee farmers and extension workers with rapid, objective disease diagnosis, supporting timely intervention and improved crop management. Ongoing research will focus on expanding dataset diversity, integrating advanced architectures, enhancing explainability, and validating the system in large-scale field deployments.

By bridging the gap between cutting-edge AI and practical agricultural needs, this project contributes to sustainable coffee production, farmer empowerment, and global food security.

## 10. Appendices

### Tables and Figures

**Table 1: Dataset Class Distribution**

Class	Number of Images
Healthy	1,179
Rust	1,023
Phoma	1,110

**Table 2: Model Architecture Summary**

Layer Name	Output Shape	Parameters
Input	(224, 224, 3)	0
MobileNetV2 Base	(7, 7, 1280)	~2.2M
GlobalAvgPooling2D	(1280)	0
Dropout (0.2)	(1280)	0
Dense (Softmax, 3)	(3)	3,843
Total		~2.2M

**Table 3: Performance Metrics**

Metric	Healthy	Rust	Phoma	Macro Avg	Weighted Avg
Precision	0.98	0.96	0.97	0.97	0.97
Recall	0.97	0.98	0.96	0.97	0.97
F1-score	0.98	0.97	0.97	0.97	0.97
Accuracy					0.97

**Table 4: Comparative Model Performance**

Model	Accuracy (%)	Model Size (MB)	Inference Time (ms)	Suitability for Mobile
MobileNetV2	97.0	14	25	Excellent
EfficientNet-B0	88.4	20	40	Good
ViT	85.1	86	120	Moderate
ResNet50	91.0	98	60	Moderate
DenseNet169	99.7	156	200	Poor